

# Hive: A data warehouse on Hadoop

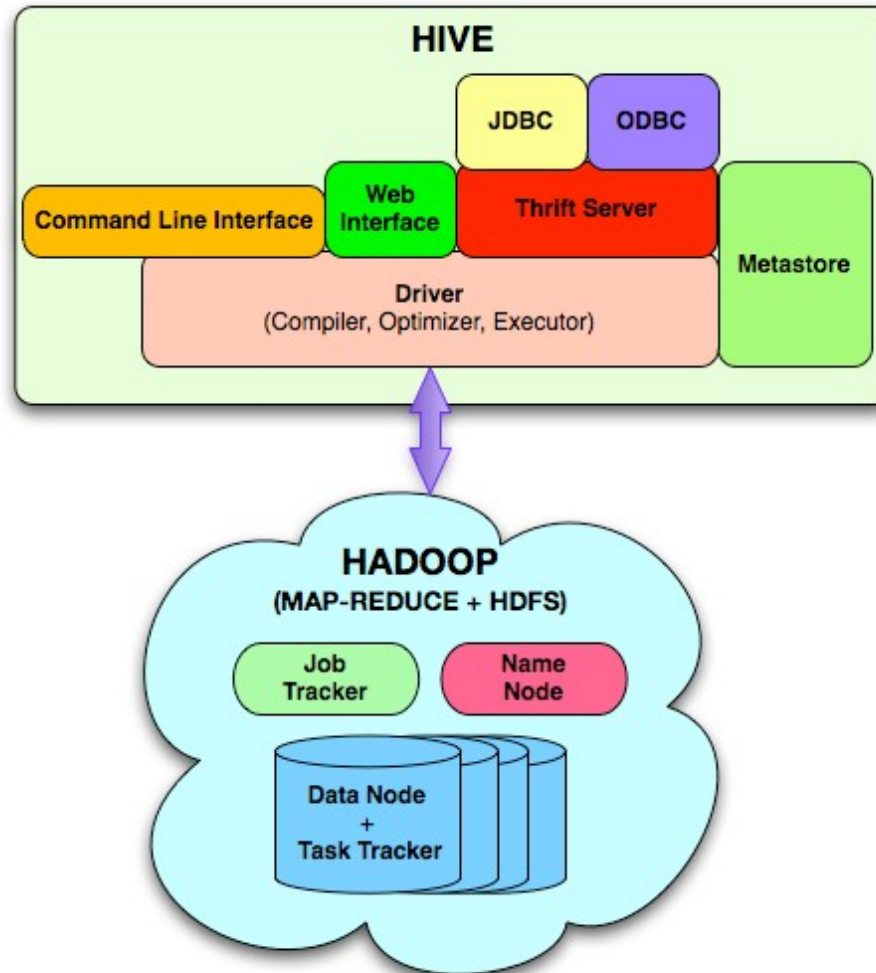
# Motivation

- Yahoo worked on Pig to facilitate application deployment on Hadoop.
  - Their need mainly was focused on unstructured data
- Simultaneously Facebook started working on deploying warehouse solutions on Hadoop that resulted in Hive.
  - The size of data being collected and analyzed in industry for business intelligence (BI) is growing rapidly making traditional warehousing solution prohibitively expensive.

# Hadoop MR

- MR is very low level and requires customers to write custom programs.
- HIVE supports queries expressed in SQL-like language called HiveQL which are compiled into MR jobs that are executed on Hadoop.
- Hive also allows MR scripts
- It also includes MetaStore that contains schemas and statistics that are useful for data explorations, query optimization and query compilation.
- At Facebook Hive warehouse contains tens of thousands of tables, stores over 700TB and is used for reporting and ad-hoc analyses by 200 Fb users.

# Hive architecture (from the paper)



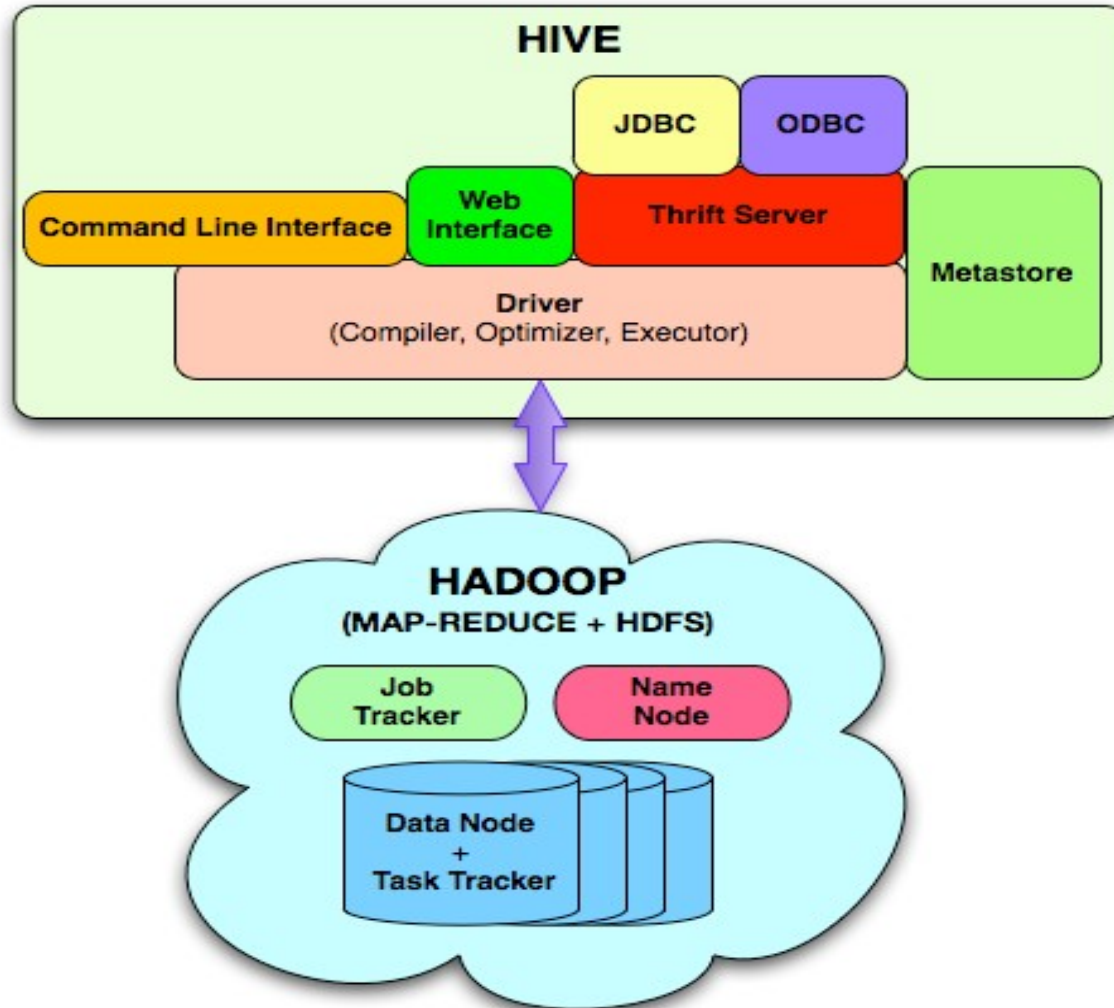
# Query Language (HiveQL)

- Subset of SQL
- Meta-data queries
- Limited equality and join predicates
- No inserts on existing tables (to preserve worm property)
  - Can overwrite an entire table


# Data Storage

- Tables are logical data units; table metadata associates the data in the table to hdfs directories.
- Hdfs namespace: tables (hdfs directory), partition (hdfs subdirectory), buckets (subdirectories within partition)
- `/user/hive/warehouse/test_table` is a hdfs directory

# System Architecture and Components



# System Architecture and Components




## Metastore

- The component that stores the system catalog and meta data about tables, columns, partitions etc.
- Stored on a traditional RDBMS




# System Architecture and Components



- **Driver**

The component that manages the lifecycle of a HiveQL statement as it moves through Hive. The driver also maintains a session handle and any session statistics.


# System Architecture and Components



- ## Query Compiler

The component that compiles HiveQL into a directed acyclic graph of map/reduce tasks.

# System Architecture and Components




## •Optimizer

consists of a chain of transformations such that the operator DAG resulting from one transformation is passed as input to the next transformation

Performs tasks like Column Pruning , Partition Pruning, Repartitioning of Data


# System Architecture and Components



- **Execution Engine**

The component that executes the tasks produced by the compiler in proper dependency order. The execution engine interacts with the underlying Hadoop instance.


# System Architecture and Components



- **HiveServer**

The component that provides a thrift interface and a JDBC/ODBC server and provides a way of integrating Hive with other applications.

# System Architecture and Components



- **Client Components**

Client component like Command Line Interface(CLI), the web UI and JDBC/ODBC driver.

# Hive Query Language

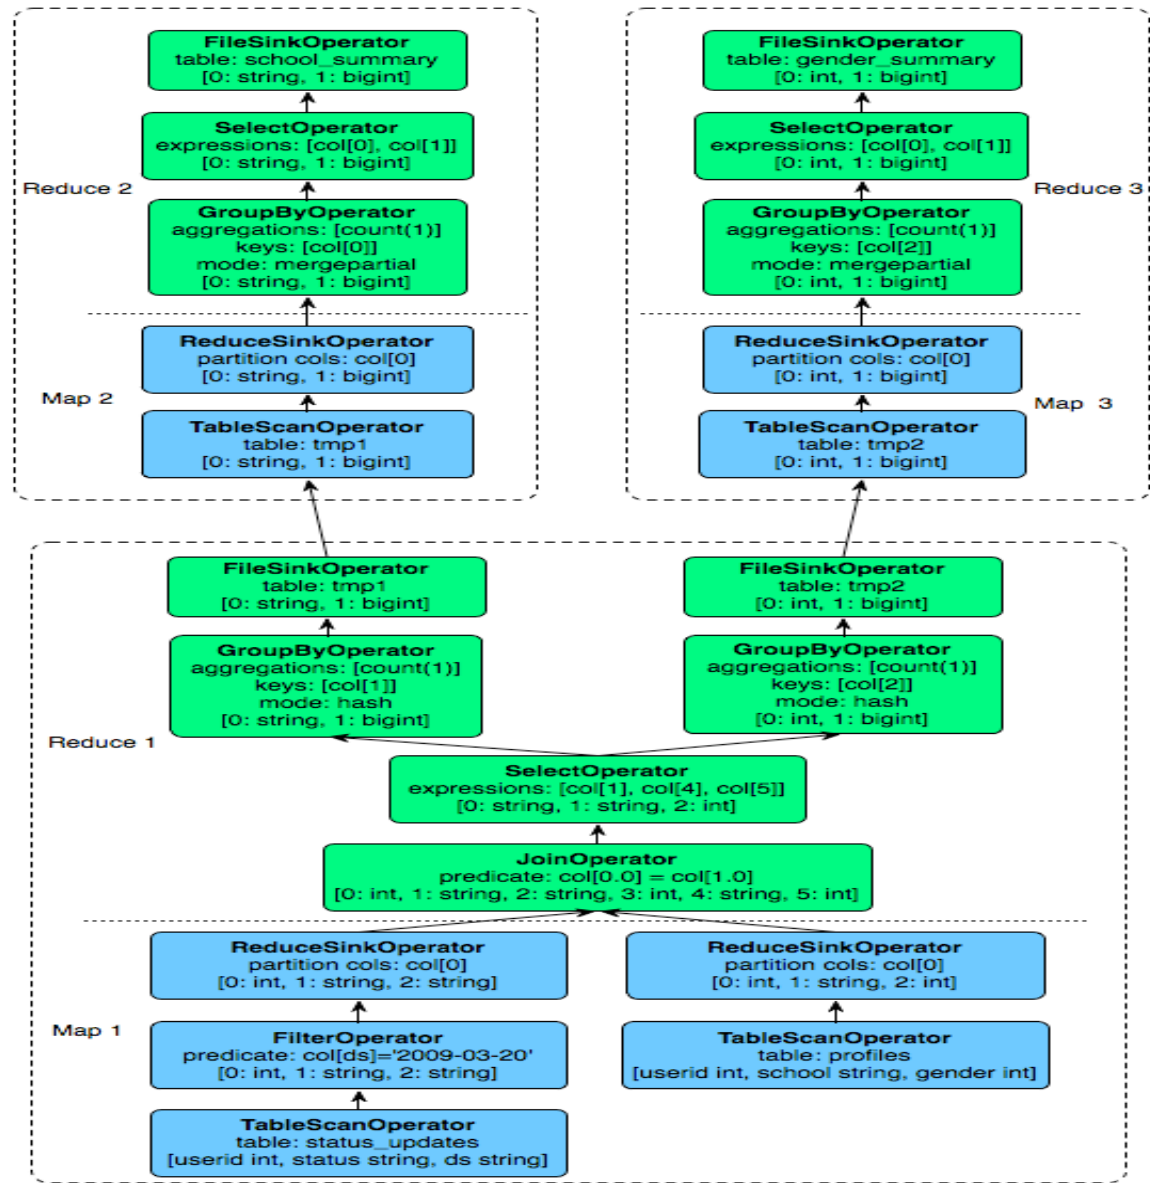
- Basic SQL
  - From clause sub-query
  - ANSI JOIN (equi-join only)
  - Multi-Table insert
  - Multi group-by
  - Sampling
  - Objects Traversal
- Extensibility
  - Pluggable Map-reduce scripts using TRANSFORM

# Architecture

- Metastore: stores system catalog
- Driver: manages life cycle of HiveQL query as it moves thru' HIVE; also manages session handle and session statistics
- Query compiler: Compiles HiveQL into a directed acyclic graph of map/reduce tasks
- Execution engines: The component executes the tasks in proper dependency order; interacts with Hadoop
- HiveServer: provides Thrift interface and JDBC/ODBC for integrating other applications.
- Client components: CLI, web interface, jdbc/odbc interface
- Extensibility interface include SerDe, User Defined Functions and User Defined Aggregate Function



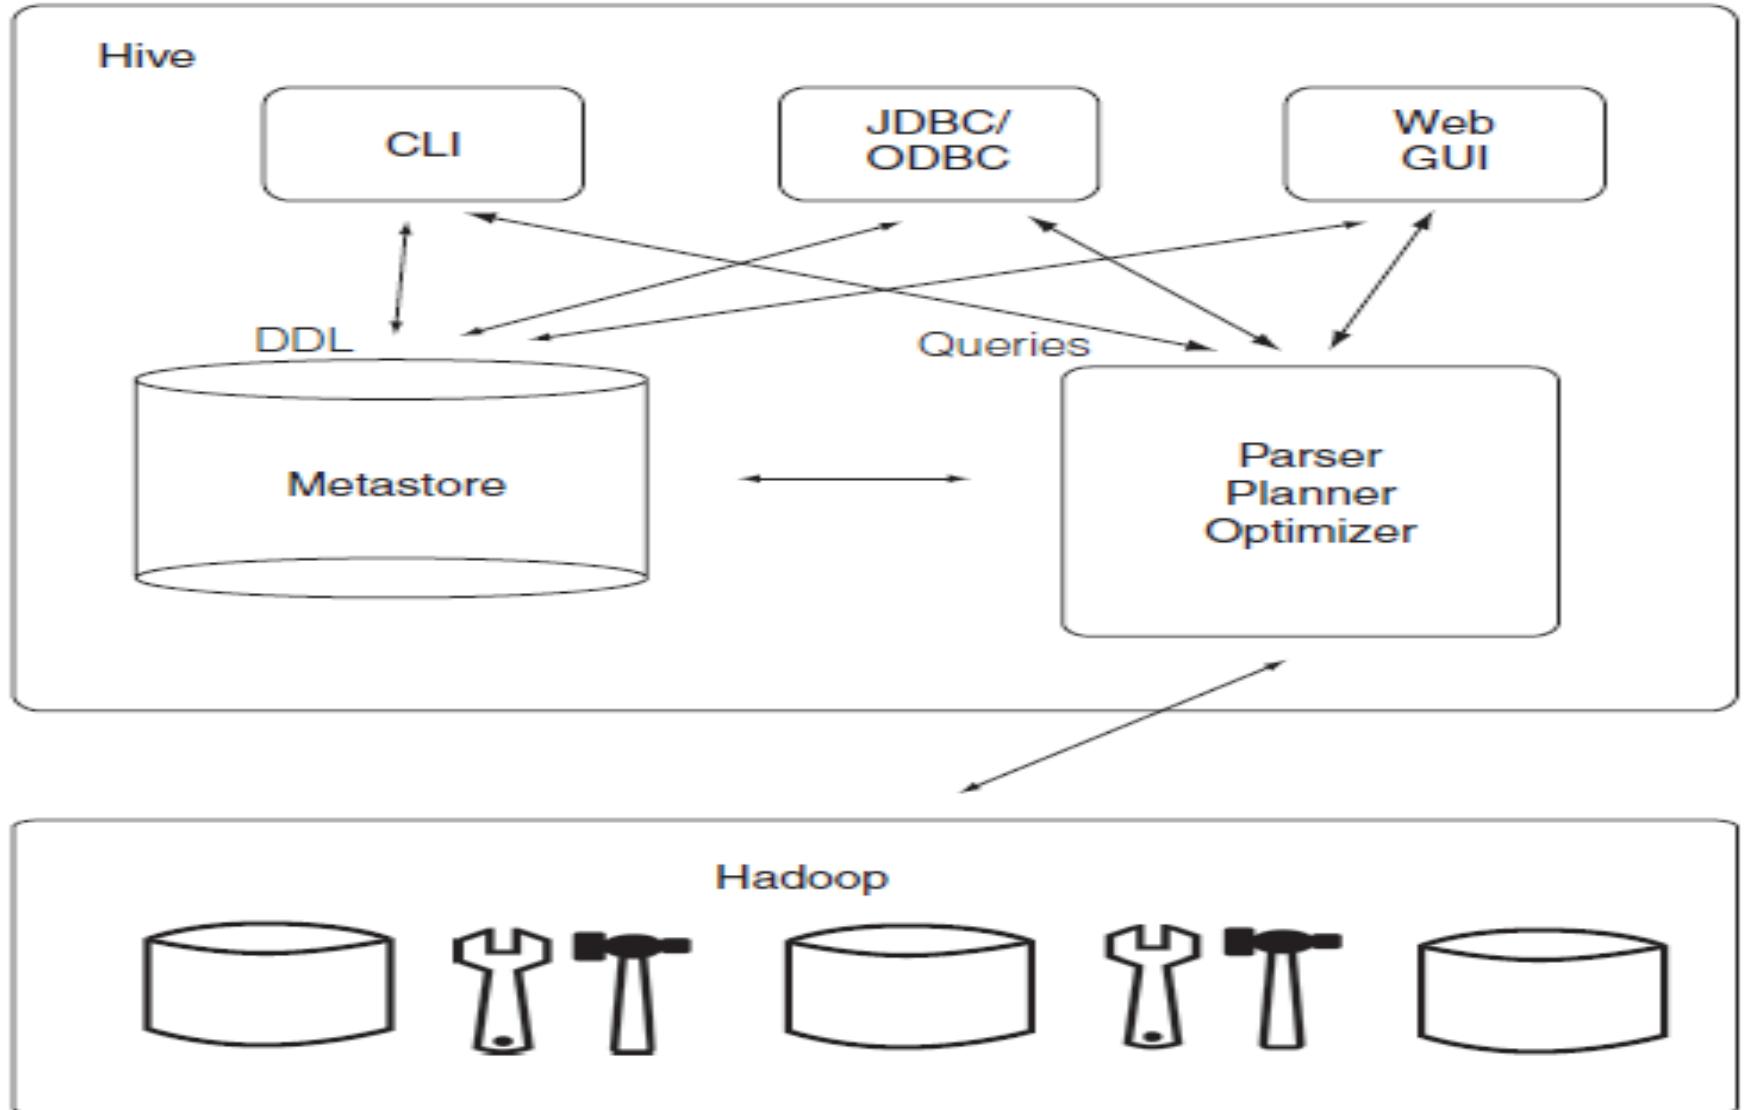
# Sample Query Plan



# Hive Usage in Facebook

- Hive and Hadoop are extensively used in Facebook for different kinds of operations.
- 700 TB = 2.1Petabyte after replication!
- Think of other application model that can leverage Hadoop MR.

# Hive Operations



# Hive Operations:tables

```
hadoop fs -mkdir /tmp
```

```
hadoop fs -mkdir /user/hive/warehouse
```

```
hadoop fs -chmod g+w /tmp
```

```
hadoop fs -chmod g+w /user/hive/warehouse
```

1. Hadoop needs to be up and running already.  
In addition, you need to set up a couple directories in HDFS for Hive to use.
- 2.It's good to let Hive manage your data if you plan on using Hive to query it.  
But if you already have your data in HDFS  
Hive can work with them too.
- 3.Hive will take your data as is and  
won't optimize.
- 4.Hive does NOT requires data to be in some special  
Hive format.

# Hive Operations: WebGUI

```
<property>
  <name>hive.hwi.war.file</name>
  <value>lib/hive-hwi-0.12.0.2.0.6.0-76.war</value>
  <description>This sets the path to the HWI war file, relative to ${HIVE_HOME}. </description>
</property>
```

```
mohit@NoMind ~/Work/BigData $ hive --service hwi
14/01/13 13:35:18 INFO hwi.HWIServer: HWI is starting up
14/01/13 13:35:18 INFO Configuration.deprecation: mapred.inpu
```

localhost:9999/hwi/index.jsp

## Hive Web Interface

Link

USER

- Home
- Authorize

DATABASE

- Browse Schema

SESSIONS

- Create Session
- List Sessions

### Hive Web Interface

The Hive Web Interface (HWI) offers an a and line interface (CLI). Once authenticated users can start a new WebSession

1. Hive with no parameters defaults to CLI  
Or explicitly fire up WebGUI

# Hive Operations

Metastore is a traditional DB and can be configured with hive-site.xml.

```
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"
<configuration>
  <property>
    <name>hive.metastore.local</name>
    <value>>false</value>
    <description>JDBC connect string for a JDBC metastore</description>
  </property>
  <property>
    <name>javax.jdo.option.ConnectionURL</name>
    <value>jdbc:derby:$HIVE_HOME/../../resources/hive_data;databaseName=metastore_db;create=true</value>
    <description>JDBC connect string for a JDBC metastore</description>
  </property>
  <property>
    <name>javax.jdo.option.ConnectionDriverName</name>
    <value>org.apache.derby.jdbc.EmbeddedDriver</value>
    <description>Driver class name for a JDBC metastore</description>
  </property>
```

# Hive Operations:creating tables

```
hive> CREATE TABLE cite (citing INT, cited INT)
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY ','
> STORED AS TEXTFILE;
```

OK

Time taken: 0.246 seconds

```
hive> SHOW TABLES;
```

OK

cite

Time taken: 0.053 seconds

```
hive> DESCRIBE cite;
```

OK

citing int

cited int

Time taken: 0.13 seconds

# Hive Operations:loading data

```
hive> LOAD DATA LOCAL INPATH 'cite75_99.txt'
> OVERWRITE INTO TABLE cite;
Copying data from file:/root/cite75_99.txt
Loading data to table cite
OK
Time taken: 9.51 seconds
```

Contents of directory [/user/hive/warehouse/cite](#)

Goto :

[Go to parent directory](#)

Name	Type	Size	Replication	Block Size	Modification Time	Permission	Owner	Group
<a href="#">cite75_99.txt</a>	file	251.84 MB	1	128 MB	2014-01-13 13:54	rw-r--r--	mohit	supergroup

[Go back to DFS home](#)

```
hive> SELECT * FROM cite LIMIT 10;
OK
NULL      NULL
3858241   956203
3858241   1324234
3858241   3398406
3858241   3557384
3858241   3634889
3858242   1515701
3858242   3319261
3858242   3668705
3858242   3707004
Time taken: 0.17 seconds
```



# Hive Operations: querying data

```
hive> SELECT COUNT(1) FROM cite;
Total MapReduce jobs = 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
    set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
    set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
    set mapred.reduce.tasks=<number>
Starting Job = job_200908250716_0001, Tracking URL = http://ip-10-244-199-143.ec2.internal:50030/jobdetails.jsp?jobid=job_200908250716_0001
Kill Command = /usr/lib/hadoop/bin/hadoop job -Dmapred.job.tracker=ip-10-244-199-143.ec2.internal:9001 -kill job_200908250716_0001
map = 0%,    reduce = 0%
map = 12%,    reduce = 0%
map = 25%,    reduce = 0%
map = 30%,    reduce = 0%
map = 34%,    reduce = 0%
map = 43%,    reduce = 0%
map = 53%,    reduce = 0%
map = 62%,    reduce = 0%
map = 71%,    reduce = 0%
map = 75%,    reduce = 0%
map = 79%,    reduce = 0%
map = 88%,    reduce = 0%
map = 97%,    reduce = 0%
map = 99%,    reduce = 0%
map = 100%,    reduce = 0%
map = 100%,    reduce = 67%
map = 100%,    reduce = 100%
Ended Job = job_200908250716_0001
OK
16522439
Time taken: 85.153 seconds
```

# Hive Operations: querying data(group by)

```
CREATE TABLE records (year STRING, temperature INT, quality INT)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY '\t';
```

```
LOAD DATA LOCAL INPATH 'input/ncdc/micro-tab/sample.txt'
OVERWRITE INTO TABLE records;
```

```
% ls /user/hive/warehouse/records/
sample.txt
```

```
hive> SELECT year, MAX(temperature)
> FROM records
> WHERE temperature != 9999
> AND (quality = 0 OR quality = 1 OR quality = 4 OR quality = 5 OR quality = 9)
> GROUP BY year;
```

1949	111
1950	22

# Hive Operations: clauses

Feature	SQL	HiveQL
Updates	UPDATE, INSERT, DELETE	INSERT OVERWRITE TABLE (populates whole table or partition)
Transactions	Supported	Not supported
Indexes	Supported	Not supported
Latency	Sub-second	Minutes
Data types	Integral, floating point, fixed point, text and binary strings, temporal	Integral, floating point, boolean, string, array, map, struct
Functions	Hundreds of built-in functions	Dozens of built-in functions
Multitable inserts	Not supported	Supported
Create table as select	Not valid SQL-92, but found in some databases	Supported
Select	SQL-92	Single table or view in the FROM clause. SORT BY for partial ordering. LIMIT to limit number of rows returned.

# Hive Operations:clauses

Feature	SQL	HiveQL
Joins	SQL-92 or variants (join tables in the FROM clause, join condition in the WHERE clause)	Inner joins, outer joins, semi joins, map joins. SQL-92 syntax, with hinting.
Subqueries	In any clause. Correlated or noncorrelated.	Only in the FROM clause. Correlated subqueries not supported
Views	Updatable. Materialized or nonmaterialized.	Read-only. Materialized views not supported
Extension points	User-defined functions. Stored procedures.	User-defined functions. MapReduce scripts.

# Hive Operations:types

Category	Type	Description	Literal examples
Primitive	TINYINT	1-byte (8-bit) signed integer, from -128 to 127	1
	SMALLINT	2-byte (16-bit) signed integer, from -32,768 to 32,767	1
	INT	4-byte (32-bit) signed integer, from -2,147,483,648 to 2,147,483,647	1
	BIGINT	8-byte (64-bit) signed integer, from -9,223,372,036,854,775,808 to 9,223,372,036,854,775,807	1
	FLOAT	4-byte (32-bit) single-precision floating-point number	1.0
	DOUBLE	8-byte (64-bit) double-precision floating-point number	1.0
	BOOLEAN	true/false value	TRUE
	STRING	Character string	'a', "a"
	BINARY	Byte array	Not supported
	TIMESTAMP	Timestamp with nanosecond precision	1325502245000, '2012-01-02 03:04:05.123456789'

# Hive Operations:types

Category	Type	Description	Literal examples
Complex	ARRAY	An ordered collection of fields. The fields must all be of the same type.	<code>array(1, 2)</code> <sup>a</sup>
	MAP	An unordered collection of key-value pairs. Keys must be primitives; values may be any type. For a particular map, the keys must be the same type, and the values must be the same type.	<code>map('a', 1, 'b', 2)</code>
	STRUCT	A collection of named fields. The fields may be of different types.	<code>struct('a', 1, 1.0)</code>

# Hive

## Operations:types(complex)

```
CREATE TABLE complex (  
  col1 ARRAY<INT>,  
  col2 MAP<STRING, INT>,  
  col3 STRUCT<a:STRING, b:INT, c:DOUBLE>  
>;
```

```
hive> SELECT col1[0], col2['b'], col3.c FROM complex;
```

```
1    2    1.0
```

# Hive Operations:partitions

```
CREATE TABLE logs (ts BIGINT, line STRING)
PARTITIONED BY (dt STRING, country STRING);
```

```
LOAD DATA LOCAL INPATH 'input/hive/partitions/file1'
INTO TABLE logs
PARTITION (dt='2001-01-01', country='GB');
```

```
/user/hive/warehouse/logs/dt=2010-01-01/country=GB/file1
                                     /file2
                                   /country=US/file3
    /dt=2010-01-02/country=GB/file4
                                   /country=US/file5
                                     /file6
```



# Hive Operations:partitions

```
hive> SHOW PARTITIONS logs;  
dt=2001-01-01/country=GB  
dt=2001-01-01/country=US  
dt=2001-01-02/country=GB  
dt=2001-01-02/country=US
```

```
SELECT ts, dt, line  
FROM logs  
WHERE country='GB';
```

# Hive Operations:buckets

```
CREATE TABLE bucketed_users (id INT, name STRING)
CLUSTERED BY (id) INTO 4 BUCKETS;
```

```
CREATE TABLE bucketed_users (id INT, name STRING)
CLUSTERED BY (id) SORTED BY (id ASC) INTO 4 BUCKETS;
```

```
hive> SELECT * FROM users;
```

```
0      Nat
2      Joe
3      Kay
4      Ann
```

```
INSERT OVERWRITE TABLE bucketed_users
SELECT * FROM users;
```

1. Bucketing imposes extra Structure on the table, Hive can exploit buckets when joining of two tables. can be efficiently implemented as a map-side join.
2. Buckets help in sampling for queries run Large data set.

# Hive Operations:buckets

```
hive> SELECT * FROM bucketed_users  
      > TABLESAMPLE(BUCKET 1 OUT OF 4 ON id);  
0      Nat  
4      Ann
```

```
hive> SELECT * FROM bucketed_users  
      > TABLESAMPLE(BUCKET 1 OUT OF 2 ON id);  
0      Nat  
4      Ann  
2      Joe
```

```
hive> SELECT * FROM users  
      > TABLESAMPLE(BUCKET 1 OUT OF 4 ON rand());  
2      Joe
```

# Hive Operations:format

```
DROP TABLE IF EXISTS compressed_users;  
CREATE TABLE compressed_users (id INT, name STRING) STORED AS SEQUENCEFILE;  
SET hive.exec.compress.output=true;  
SET mapred.output.compress=true;  
SET mapred.output.compression.codec=org.apache.hadoop.io.compress.GzipCodec;  
INSERT OVERWRITE TABLE compressed_users SELECT * FROM users;
```

# Hive Operations:format(RC)

**Logical table**

	col1	col2	col3
row1	1	2	3
row2	4	5	6
row3	7	8	9
row4	10	11	12

**Row-oriented layout (SequenceFile)**

row1	row2	row3	row4
1 2 3	4 5 6	7 8 9	10 11 12

**Column-oriented layout (RCFile)**

row split 1						row split 2					
col1		col2		col3		col1		col2		col3	
1	4	2	5	3	6	7	10	8	11	9	12

# Hive Operations:joins

```
hive> SELECT * FROM sales;
```

```
Joe      2
```

```
Hank     4
```

```
Ali      0
```

```
Eve      3
```

```
Hank     2
```

```
hive> SELECT * FROM things;
```

```
2      Tie
```

```
4      Coat
```

```
3      Hat
```

```
1      Scarf
```

# Hive Operations:joins

```
hive> SELECT sales.*, things.*  
      > FROM sales JOIN things ON (sales.id = things.id);
```

Joe	2	2	Tie
Hank	2	2	Tie
Eve	3	3	Hat
Hank	4	4	Coat

```
hive> SELECT sales.*, things.*  
      > FROM sales LEFT OUTER JOIN things ON (sales.id = things.id);
```

Ali	0	NULL	NULL
Joe	2	2	Tie
Hank	2	2	Tie
Eve	3	3	Hat
Hank	4	4	Coat

# Hive Operations:joins

```
hive> SELECT sales.*, things.*  
      > FROM sales RIGHT OUTER JOIN things ON (sales.id = things.id);
```

NULL	NULL	1	Scarf
Joe	2	2	Tie
Hank	2	2	Tie
Eve	3	3	Hat
Hank	4	4	Coat

```
hive> SELECT sales.*, things.*  
      > FROM sales FULL OUTER JOIN things ON (sales.id = things.id);
```

Ali	0	NULL	NULL
NULL	NULL	1	Scarf
Joe	2	2	Tie
Hank	2	2	Tie
Eve	3	3	Hat
Hank	4	4	Coat



# Hive Operations:joins

```
SELECT *  
FROM things  
WHERE things.id IN (SELECT id from sales);
```

```
hive> SELECT *  
      > FROM things LEFT SEMI JOIN sales ON (sales.id = things.id);  
2    Tie
```

```
SELECT /*+ MAPJOIN(things) */ sales.*, things.*  
FROM sales JOIN things ON (sales.id = things.id);
```

# Hive

## Operations:subqueries(from )

```
SELECT station, year, AVG(max_temperature)
FROM (
  SELECT station, year, MAX(temperature) AS max_temperature
  FROM records2
  WHERE temperature != 9999
  AND (quality = 0 OR quality = 1 OR quality = 4 OR quality = 5 OR quality = 9)
  GROUP BY station, year
```

# Hive Operations:views

```
CREATE VIEW valid_records
AS
SELECT *
FROM records2
WHERE temperature != 9999
    AND (quality = 0 OR quality = 1 OR quality = 4 OR quality = 5 OR quality = 9);
```

```
CREATE VIEW max_temperatures (station, year, max_temperature)
AS
SELECT station, year, MAX(temperature)
FROM valid_records
GROUP BY station, year;
```

# Hive Operations:serde

```
CREATE TABLE stations (usaf STRING, wban STRING, name STRING)
ROW FORMAT SERDE 'org.apache.hadoop.hive.contrib.serde2.RegexSerDe'
WITH SERDEPROPERTIES (
  "input.regex" = "(\d{6}) (\d{5}) (.{29}) .*"
);
```

```
LOAD DATA LOCAL INPATH "input/ncdc/metadata/stations-fixed-width.txt"
INTO TABLE stations;
```

```
hive> SELECT * FROM stations LIMIT 4;
010000      99999      BOGUS NORWAY
010003      99999      BOGUS NORWAY
010010      99999      JAN MAYEN
010013      99999      ROST
```

# Hive Operations:serde(sort)

```
hive> FROM records2
      > SELECT year, temperature
      > DISTRIBUTE BY year
      > SORT BY year ASC, temperature DESC;
1949      111
1949      78
1950      22
1950       0
1950     -11
```

# Hive Operations:serde

```
hive> CREATE EXTERNAL TABLE logs_20120101 (
    host STRING,
    identity STRING,
    user STRING,
    time STRING,
    request STRING,
    status STRING,
    size STRING)
ROW FORMAT SERDE 'org.apache.hadoop.hive.contrib.serde2.RegexSerDe'
WITH SERDEPROPERTIES (
    "input.regex" =
        "([^\ ]*) ([^\ ]*) ([^\ ]*) (-|\\[[^\]]*\])"
        "([^\\" ]*|\"[^\"]*\\") (-|[0-9]*) (-|[0-9]*)",
    "output.format.string"="%1$s %2$s %3$s %4$s %5$s %6$s %7$s"
)
STORED AS TEXTFILE LOCATION '/data/logs/20120101/';
```

The regular expression used to match and extract groups that are mapped to the table columns. Also note that there's a single space separator where the regular expression is split across two lines.


- Determines the order and formatting of the table when it's being written.

```
"input.regex" =  
"([ ^ ]*) ([^ ]*) ([^ ]*) (-|\\[[^\\]]*\\]) ([^ \"]*|\"[^\"]*\" ) (-|[0-9]*) (-|[0-9]*)",  
      ↑       ↑       ↑       ↑       ↑       ↑  
      space   space   space   space   space   space
```

# Hive Operations:serde

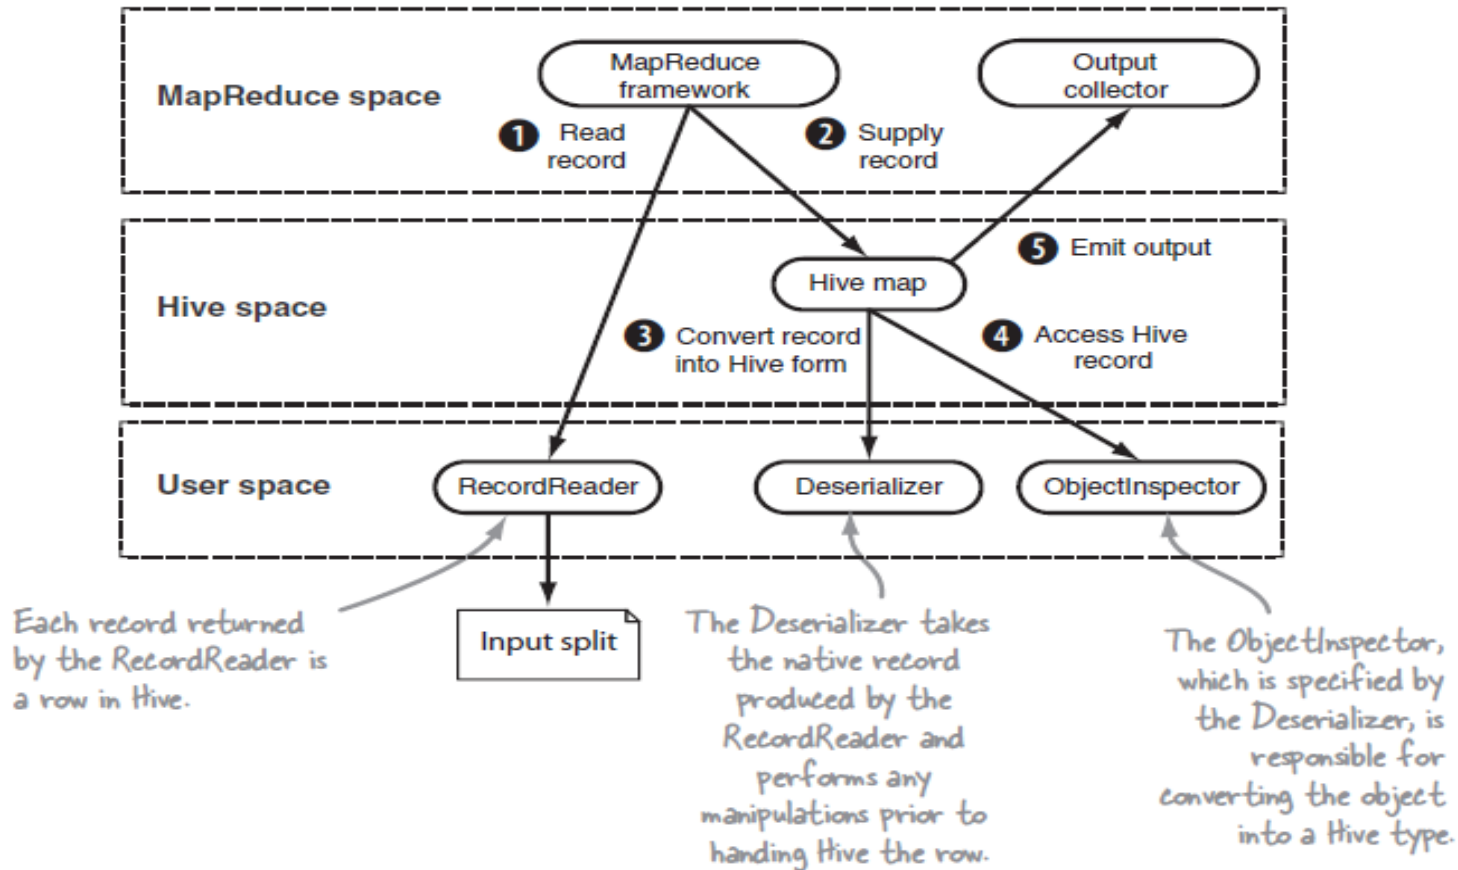
```
hive> add jar $HIVE_HOME/lib/hive-contrib-0.7.1-cdh3u2.jar;  
hive> SELECT host, request FROM logs_20120101 LIMIT 10;
```

```
89.151.85.133  "GET /movie/127Hours HTTP/1.1"  
212.76.137.2   "GET /movie/BlackSwan HTTP/1.1"  
74.125.113.104 "GET /movie/TheFighter HTTP/1.1"  
212.76.137.2   "GET /movie/Inception HTTP/1.1"  
127.0.0.1      "GET /movie/TrueGrit HTTP/1.1"  
10.0.12.1      "GET /movie/WintersBone HTTP/1.1"
```



CDH Hive is installed under /usr/lib/hive. You should substitute `$HIVE_HOME` with the location of your Hive installation because Hive doesn't expand environment variables.

# Hive Operations:serde




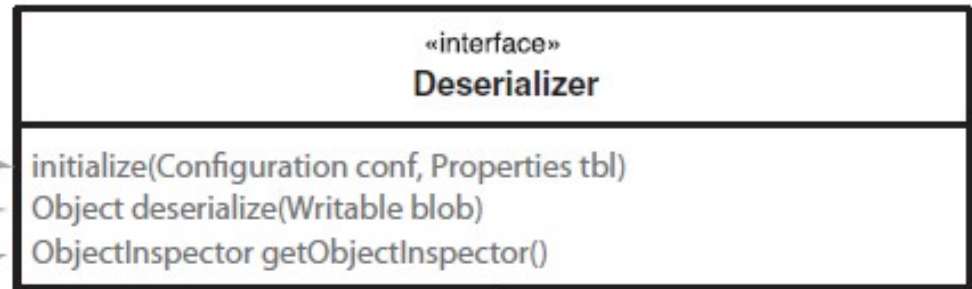


# Hive Operations:serde

Validate that the user-specified fields conform to the specifications of the implementation.

Convert the object created by the InputFormat into a Hive form.

Create an ObjectInspector, which is a class that knows how to access individual fields within the Object returned by the deserialize method.



@Override

```
public void initialize(Configuration conf, Properties tbl)
```

```
    throws SerDeException {
```

```
        inputRegex = tbl.getProperty("input.regex");
```

Read the regular expression  
from the table definition.

```
        String columnNameProperty = tbl.getProperty(  
            Constants.LIST_COLUMNS);
```

Read the column names from the table definition.

```
        String columnTypeProperty = tbl.getProperty(  
            Constants.LIST_COLUMN_TYPES);
```

Read the column types from  
the table definition.

```
        inputPattern = Pattern.compile(inputRegex, ...);
```

Construct the Java Pattern object, which  
will be used in the deserialize method.

```
        List<String> columnNames = Arrays.asList(  
            columnNameProperty.split(","));
```

Tokenize the column names.

```
        List<TypeInfo> columnTypes = TypeInfoUtils  
            .getTypeInfosFromTypeString(columnTypeProperty);
```

Tokenize the column types.

Ensure that each column  
type is a String.

```
        for (int c = 0; c < numColumns; c++) {  
            if (!columnTypes.get(c).equals(TypeInfoFactory.stringTypeInfo)) {  
                throw new SerDeException(...);  
            }  
        }  
    }
```

```
        List<ObjectInspector> columnOIs = new ArrayList<ObjectInspector>(  
            columnNames.size());
```

# Hive Operations:serde


```
for (int c = 0; c < numColumns; c++) {  
    columnOIs.add(  
        PrimitiveObjectInspectorFactory.javaStringObjectInspector);  
}  
  
rowOI =  
    ObjectInspectorFactory.getStandardStructObjectInspector(  
        columnNames, columnOIs);  
}  
  
@Override  
public ObjectInspector getObjectInspector() throws SerDeException {  
    return rowOI;  
}
```

Create a primitive *ObjectInspector* for each field.


Create a structure *ObjectInspector* for the row that works with List and Array-based representations of rows.




# Hive Operations:serde


```
@Override
public Object deserialize(Writable blob) throws SerDeException {

    Text rowText = (Text) blob;  Convert the Writable into a Text object.

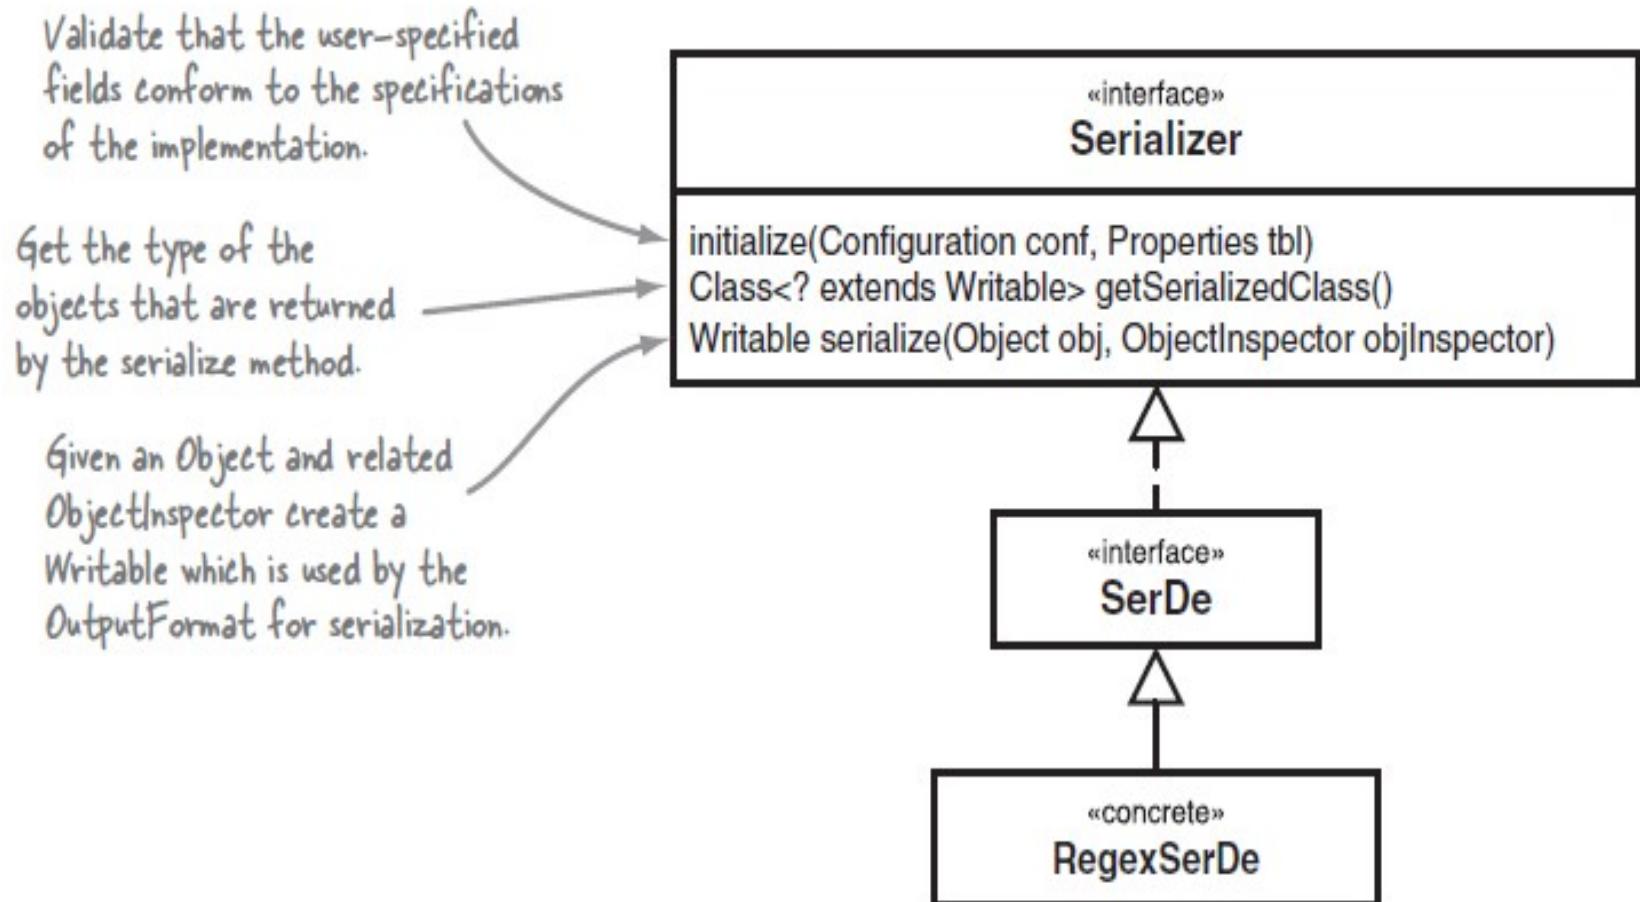
    Matcher m = inputPattern.matcher(rowText.toString());

    // If do not match, ignore the line, return a row with all nulls.
    if (!m.matches()) {
        return null;  If the regular expression didn't match the record, return null.
    }

    // Otherwise, return the row.
    for (int c = 0; c < numColumns; c++) {
        try {
            row.set(c, m.group(c + 1));  For each group in the regular expression, set the appropriate column
        } catch (RuntimeException e) {  in the array—the row is a reusable ArrayList that was created in
            row.set(c, null);  the initialize method, which was omitted for brevity.
        }
    }
    return row;
}
```

 Set a null if you ran out of groups.

# Hive Operations:serde



```
@Override
public Class<? extends Writable> getSerializedClass() {
    return Text.class;
}
```

← Tell Hive that your serialize method produces Text objects.

```
Object[] outputFields;
Text outputRowText;
```

```
@Override
public Writable serialize(Object obj, ObjectInspector objInspector)
    throws SerDeException {
```

```
    StructObjectInspector outputRowOI =
        (StructObjectInspector) objInspector;
    List<? extends StructField> outputFieldRefs = outputRowOI
        .getAllStructFieldRefs();
```

```
    for (int c = 0; c < numColumns; c++) {
```

```
        Object field = outputRowOI
            .getStructFieldData(obj, outputFieldRefs.get(c));
```

```
        ObjectInspector fieldOI = outputFieldRefs.get(c)
            .getFieldObjectInspector();
```

```
        StringObjectInspector fieldStringOI = (StringObjectInspector)
            fieldOI;
```

```
        outputFields[c] =
            fieldStringOI.getPrimitiveJavaObject(field);
```

```
    }
```

```
    String outputRowString = String.format(
        outputFormatString, outputFields);
```

```
    outputRowText.set(outputRowString);
    return outputRowText;
}
```

← Extract the individual  
ObjectInspector for each  
field in the table.

← Use the ObjectInspector to extract the column.

← Create the output line using all the columns with  
the format defined in the outputFormatString,  
which is set in the SerDe properties



# Hive:data

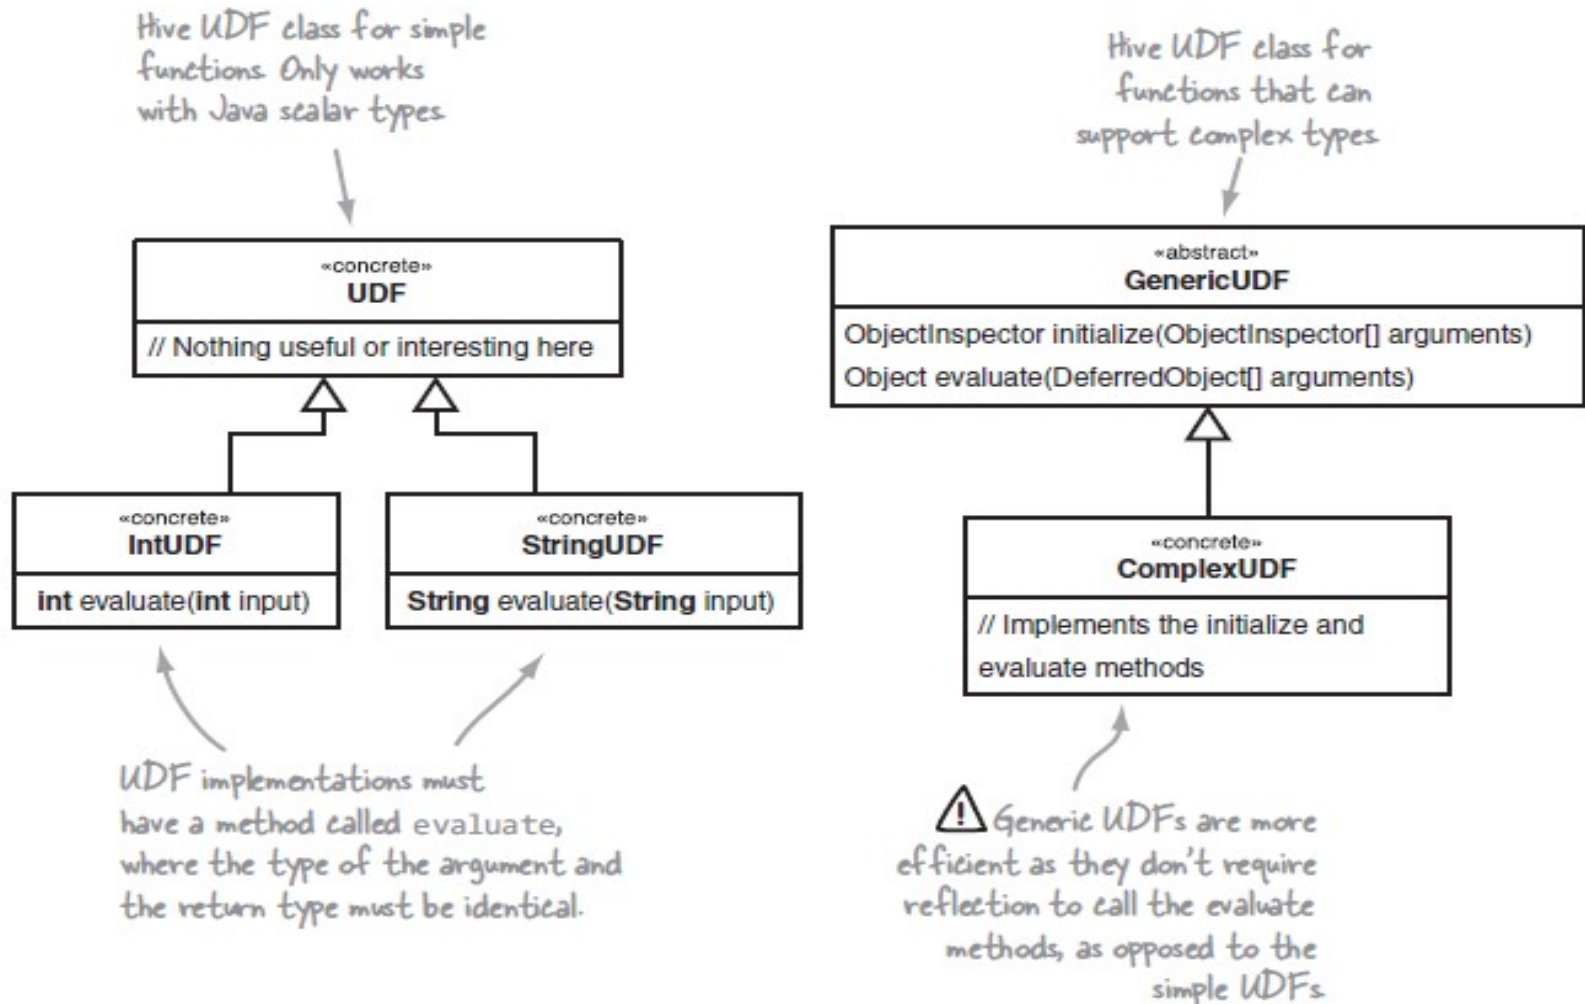
## analytics:example-1

```
> select * from logs_20120101;
OK
89.151.85.133 - - [23/Jun/2009:10:39:11 +0300] "GET /movie/127Hours HTTP/1.1" 200 766
212.76.137.2 - - [23/Jun/2009:10:39:11 +0300] "GET /movie/BlackSwan HTTP/1.1" 200 766
74.125.113.104 - - [23/Jun/2009:10:39:11 +0300] "GET /movie/TheFighter HTTP/1.1" 200 766
212.76.137.2 - - [23/Jun/2009:10:39:11 +0300] "GET /movie/Inception HTTP/1.1" 200 766
127.0.0.1 - - [23/Jun/2009:10:39:11 +0300] "GET /movie/TrueGrit HTTP/1.1" 200 766
10.0.12.1 - - [23/Jun/2009:10:39:11 +0300] "GET /movie/WintersBone HTTP/1.1" 200 766
Time taken: 0.128 seconds, Fetched: 6 row(s)
```

```
hive> select * from movie_categories;
FAILED: SemanticException [Error 10001]: Line 1:14 Table 'movie_categories' does not exist
hive> select * from movie_categories;
OK
127Hours Adventure
127Hours Drama
BlackSwan Drama
BlackSwan Thriller
TheFighter Drama
TheFighter Drama
13Assassins Action
Time taken: 0.088 seconds, Fetched: 7 row(s)
```

1. join table on movie to view categories
2. See geographical location the request is coming from
3. Understand UDFs and UDAFs
4. See dynamic and static partitioning in progress

# Hive: data analytics: example-1: UDFs and UDAFs





# Hive:data

## analytics:example-1:UDFsand UDAFs

```
public class ExtractMovieUDF extends UDF {  
    private Text result = new Text();  
    public Text evaluate(final Text t) {  
        if (t == null) { return null; }  
        String s = t.toString();  
        String[] parts = StringUtils.split(s, " ");  
        if(parts.length != 3) {  
            return null;  
        }  
        String path = parts[1];  
  
        if(!path.startsWith("/movie/")) {  
            return null;  
        }  
        result.set(path.substring(7));  
        return result;  
    }  
}
```

Looking at the UDF class, there aren't actually any methods to override to implement the function. Hive actually uses reflection to find methods whose names are evaluate and matches the arguments used in the HiveQL function call. Hive can work with both the Hadoop Writables and the Java primitives, but it's recommended to work with the Writables since they can be reused.

This UDF works on the request field, which you split into three parts: the HTTP method, the resource (which is the URL path), and the protocol.

Ignore URLs that don't pertain to the movie part of your website.

Extract the text after the leading movie path, which contains your movie title.

The Description annotation is used to provide usage information in the Hive shell (you'll see how this works following this code).

```
@Description(  
    name = "country",  
    value = "_FUNC_(ip, geolocfile) - Returns the geolocated " +  
    "country code for the IP"
```

The geolocation  
lookup class

```
)  
public class GeolocUDF extends GenericUDF {  
    private LookupService geoloc;  
    private ObjectInspectorConverters.Converter[] converters;
```

Converters, which you'll use to convert  
the input types to the types you want  
to operate with.

```
@Override  
public ObjectInspector initialize(ObjectInspector[] arguments) {  
    converters =
```

```
        new ObjectInspectorConverters.Converter[arguments.length];  
    for (int i = 0; i < arguments.length; i++) {  
        converters[i] =  
            ObjectInspectorConverters.getConverter(arguments[i],  
            PrimitiveObjectInspectorFactory.javaStringObjectInspector);
```

Create a converter that you can use in the evaluate  
method to convert all the arguments  
(which in your case are the  
IP address and geolocation file)  
from their native type into Java Strings.

```
    }  
    return PrimitiveObjectInspectorFactory  
        .getPrimitiveJavaObjectInspector(  
            PrimitiveObjectInspector.PrimitiveCategory.STRING);
```

Specify that the return type for the UDF (in  
other words the evaluate function) will be a  
Java String.

```
}
```

```
@Override
public Object evaluate(GenericUDF.DeferredObject[] arguments) {

    Text ip = (Text) converters[0].convert(arguments[0].get());
    Text filename = (Text) converters[1].convert(arguments[1].get());
```

```
    return lookup(ip, filename);
}
```

← After retrieving the IP address and geolocation filename from the arguments, call a function to perform the geolocation.

```
protected String lookup(Text ip, Text filename)
    throws HiveException {
```

```
    try {
        if (geoloc == null) {
            URL u = getClass().getClassLoader()
                .getResource(filename.toString());
            geoloc =
                new LookupService(u.getFile(),
                                LookupService.GEOIP_MEMORY_CACHE);
        }
```

← Load the geolocation data file from the distributed cache.

← Create an instance of the MaxMind Lookup class.

```
        String countryCode =
            geoloc.getCountry(ip.toString()).getCode();
```

← Perform the geolocation and extract the country code.

```
        if ("--".equals(countryCode)) {
            return null;
        }
```

```
        return countryCode;
    } catch (IOException e) {
        throw new HiveException("Caught IO exception", e);
    }
```

← Return the country code.

Dynamic partitions needs to be explicitly enabled in Hive.

hive> SET hive.exec.dynamic.partition=true;

hive> SET hive.enforce.bucketing = true;

Earlier you specified that the viewed\_movies was a bucketed table with 64 buckets. Bucketed tables are optimized for sampling because without them extracting a sample from a table requires a full table scan. Whenever you write to a bucketed table, you need to make sure that you either set hive.enforce.bucketing to true, or set mapred.reduce.tasks to the number of buckets.

hive> add jar /home/mohit/Work/BigData/resources/hive/lib/hive-contrib-0.11.0.2.0.5.0-67.jar;

Add the geolocation data file into the distributed cache.

hive> ADD jar /home/mohit/Work/BigData/TOOLS/Hive/target/Hive-1.0-SNAPSHOT-jar-with-dependencies.jar;

Add the JAR containing your UDF so that it can be used in MapReduce.

hive> ADD file /home/mohit/Work/BigData/input/hive/serde/GeoIP.dat;

hive> CREATE temporary function country\_udf AS 'org.mk.training.hive.serde.udf.GeolocUDF';

Define country\_udf as the alias for your geolocation UDF and specify the class name.

hive> CREATE temporary function movie\_udf AS 'org.mk.training.hive.serde.udf.ExtractMovieUDF';

Define an alias for your movie UDF, which extracts the movie name from the URL path.



Enable compression  
for MapReduce job  
outputs.

hive> SET hive.exec.compress.output=true;

hive> SET hive.exec.compress.intermediate = true;

Enable compression for  
intermediate map outputs.

hive> SET mapred.output.compression.codec =  
org.apache.hadoop.io.compress.SnappyCodec;

Use the Snappy  
compression codec.

hive> INSERT OVERWRITE TABLE viewed\_movies  
PARTITION (dt='2012-01-01', country)  
SELECT host, movie\_udf(request),  
country\_udf(host, "GeoIP.dat")  
FROM logs\_20120101;

An example of both a static (the dt column) and  
dynamic (the country column) partitions in action.

Call your UDF specifying the field on which it should operate  
(the host column from the logs table), and the filename of the  
geolocation data file, which is in the distributed cache.

hive> SELECT \* from viewed\_movies;

OK

89.151.85.133	127Hours	2012-01-01	GB
212.76.137.2	BlackSwan	2012-01-01	RU
212.76.137.2	Inception	2012-01-01	RU
74.125.113.104	TheFighter	2012-01-01	US
127.0.0.1	TrueGrit	2012-01-01	

The `__HIVE_DEFAULT_PARTITION__` is used  
to store records whose dynamic partition column  
value is NULL, or the empty string.

<code>__HIVE_DEFAULT_PARTITION__</code>	<code>__HIVE_DEFAULT_PARTITION__</code>	<code>__HIVE_DEFAULT_PARTITION__</code>	<code>__HIVE_DEFAULT_PARTITION__</code>
10.0.12.1	WintersBone	2012-01-01	<code>__HIVE_DEFAULT_PARTITION__</code>

# Hive:data

## analytics:example-1:partitioning and bucketing

Table directory.

Since we had two partitions we have two levels of directories that model the partitions

```
$ fs -lsr /user/hive/warehouse
/user/hive/warehouse/viewed_movies
/user/hive/warehouse/viewed_movies/dt=2012-01-01
/user/hive/warehouse/viewed_movies/dt=2012-01-01/country=GB
/user/hive/warehouse/viewed_movies/dt=2012-01-01/country=GB/000000_0
/user/hive/warehouse/viewed_movies/dt=2012-01-01/country=GB/000001_0
/user/hive/warehouse/viewed_movies/dt=2012-01-01/country=RU
/user/hive/warehouse/viewed_movies/dt=2012-01-01/country=RU/000000_0
/user/hive/warehouse/viewed_movies/dt=2012-01-01/country=RU/000001_0
/user/hive/warehouse/viewed_movies/dt=2012-01-01/country=US
/user/hive/warehouse/viewed_movies/dt=2012-01-01/country=US/000000_0
/user/hive/warehouse/viewed_movies/dt=2012-01-01/country=US/000001_0
/user/hive/warehouse/viewed_movies/dt=2012-01-01/country=__HIVE_DEFAULT_PARTITION__
/user/hive/warehouse/viewed_movies/dt=2012-01-01/country=__HIVE_DEFAULT_PARTITION__/000000_0
/user/hive/warehouse/viewed_movies/dt=2012-01-01/country=__HIVE_DEFAULT_PARTITION__/000001_0
```


At the leaf directory we have a file for each bucket number that we specified when creating the table

# Hive:data analytics:example-1:joins

```
hive> SELECT * FROM viewed_movies;
89.151.85.133    127Hours    2012-01-01    GB
212.76.137.2    BlackSwan    2012-01-01    RU
212.76.137.2    Inception    2012-01-01    RU
74.125.113.104  TheFighter    2012-01-01    US
127.0.0.1        TrueGrit     2012-01-01    __HIVE_DEFAULT_PARTITION__
10.0.12.1        WintersBone  2012-01-01    __HIVE_DEFAULT_PARTITION__
```

```
hive> SELECT * from movie_categories;
127Hours    Adventure
127Hours    Drama
BlackSwan    Drama
BlackSwan    Thriller
TheFighter    Drama
13Assassins  Action
```

```
hive> SET mapred.reduce.tasks=2;  Set the number of reducers for the job.
```

```
hive> SELECT viewed_movies.movie, movie_categories.category
FROM viewed_movies
JOIN movie_categories ON  (viewed_movies.movie = movie_categories.title);
127Hours    Adventure
127Hours    Drama
BlackSwan    Drama
BlackSwan    Thriller
TheFighter    Drama
```

Indicate that you're joining the viewed\_movies table with the movie\_categories table, and specify the columns that should be joined (viewed\_movies.movie and movie\_categories.title).

# Hive:data analytics:example-1:joins

```
hive> SELECT viewed_movies.movie, movie_categories.category
FROM viewed_movies
LEFT OUTER JOIN movie_categories ON
    (viewed_movies.movie = movie_categories.title);
```

A left outer join includes all rows from the viewed\_movies table regardless of whether or not they have a matching row in the movie\_categories table.

127Hours	Adventure
127Hours	Drama
BlackSwan	Drama
BlackSwan	Thriller
Inception	NULL
TheFighter	Drama
TrueGrit	NULL
WintersBone	NULL

You have some movies that don't have categories, in which case with a left outer join they'll have a NULL value for the category field.



# Hive:data analytics:example-1:joins

```
hive> SELECT viewed_movies.movie, movie_categories.category,  
            movie_categories.title  
FROM viewed_movies  
RIGHT OUTER JOIN movie_categories ON  
    (viewed_movies.movie = movie_categories.title);
```

NULL	Action	13Assassins
BlackSwan	Drama	BlackSwan
BlackSwan	Thriller	BlackSwan
TheFighter	Drama	TheFighter
TheFighter	Drama	TheFighter
127Hours	Adventure	127Hours
127Hours	Drama	127Hours

A right outer join includes all rows from the `movie_categories` tables, and only matching rows from the `viewed_movies` table.

Because all the rows from the `movie_categories` are included, rows that don't have a matching entry in the `viewed_movies` table will contain a `NULL` value for any columns from that table.

# Hive:data analytics:example-1:joins

```
hive> SELECT viewed_movies.movie, movie_categories.category,  
            movie_categories.title  
FROM viewed_movies  
FULL OUTER JOIN movie_categories ON  
    (viewed_movies.movie = movie_categories.title);
```

NULL	Action	13Assassins
BlackSwan	Drama	BlackSwan
BlackSwan	Thriller	BlackSwan
TheFighter	Drama	TheFighter
TheFighter	Drama	TheFighter
TrueGrit	NULL	NULL
WintersBone	NULL	NULL
127Hours	Adventure	127Hours
127Hours	Drama	127Hours
Inception	NULL	NULL

A full outer join includes all rows from both tables regardless of whether there is a match.

Just like with the left and right outer joins, any rows that fail to match will contain *NULL* values for the table with no corresponding entry.

# Hive:data analytics:example-1:joins

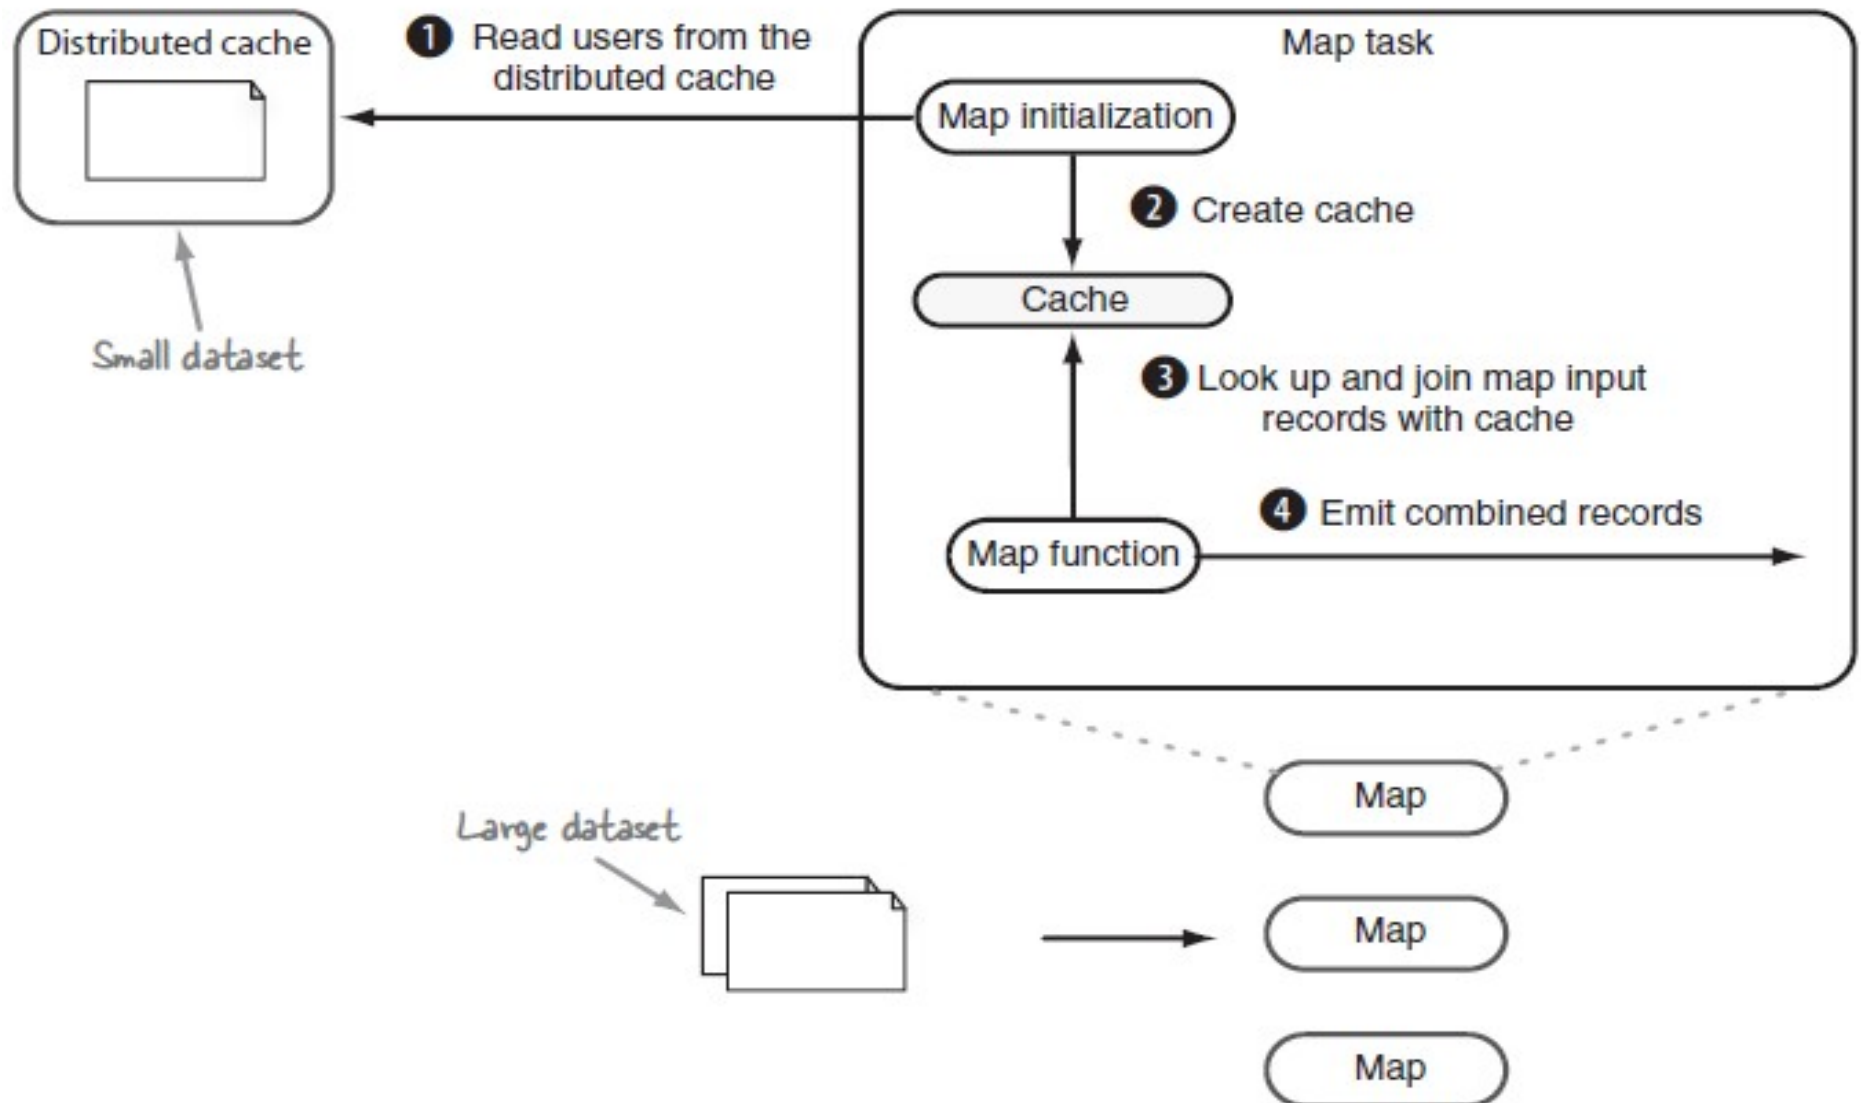
```
hive> SELECT /*+ MAPJOIN(movie_categories) */  
        viewed_movies.movie, movie_categories.category  
FROM viewed_movies  
JOIN movie_categories  
ON viewed_movies.movie = movie_categories.title;
```

The hint which triggers the map join and also tells  
Hive which table (movie\_categories) to cache.

The files in the small tables are smaller than the value specified in `hive.mapjoin.smalltable.filesize`, whose default value is set at 25 MB.

The memory utilization of loading the small table must be less than `hive.mapjoin.localtask.max.memory.usage`, which is set at 0.90 (90 percent) by default.

# Hive: data analytics: example-1: joins



# Hive:data analytics:example-1:joins

```
hive> SET hive.auto.convert.join = true;
```

Enable Hive's automatic join optimization to convert repartition joins to replicated joins if one of the tables is small enough.

```
hive> SELECT viewed_movies.movie, movie_categories.category
FROM viewed_movies
JOIN movie_categories
ON viewed_movies.movie = movie_categories.title;
```

```
...
Mapred Local Task Succeeded. Convert the Join into MapJoin
...
```

This output tells you that Hive decided to go ahead and use a replicated join.


# Hive:data

## analytics:example-1:semi-joins

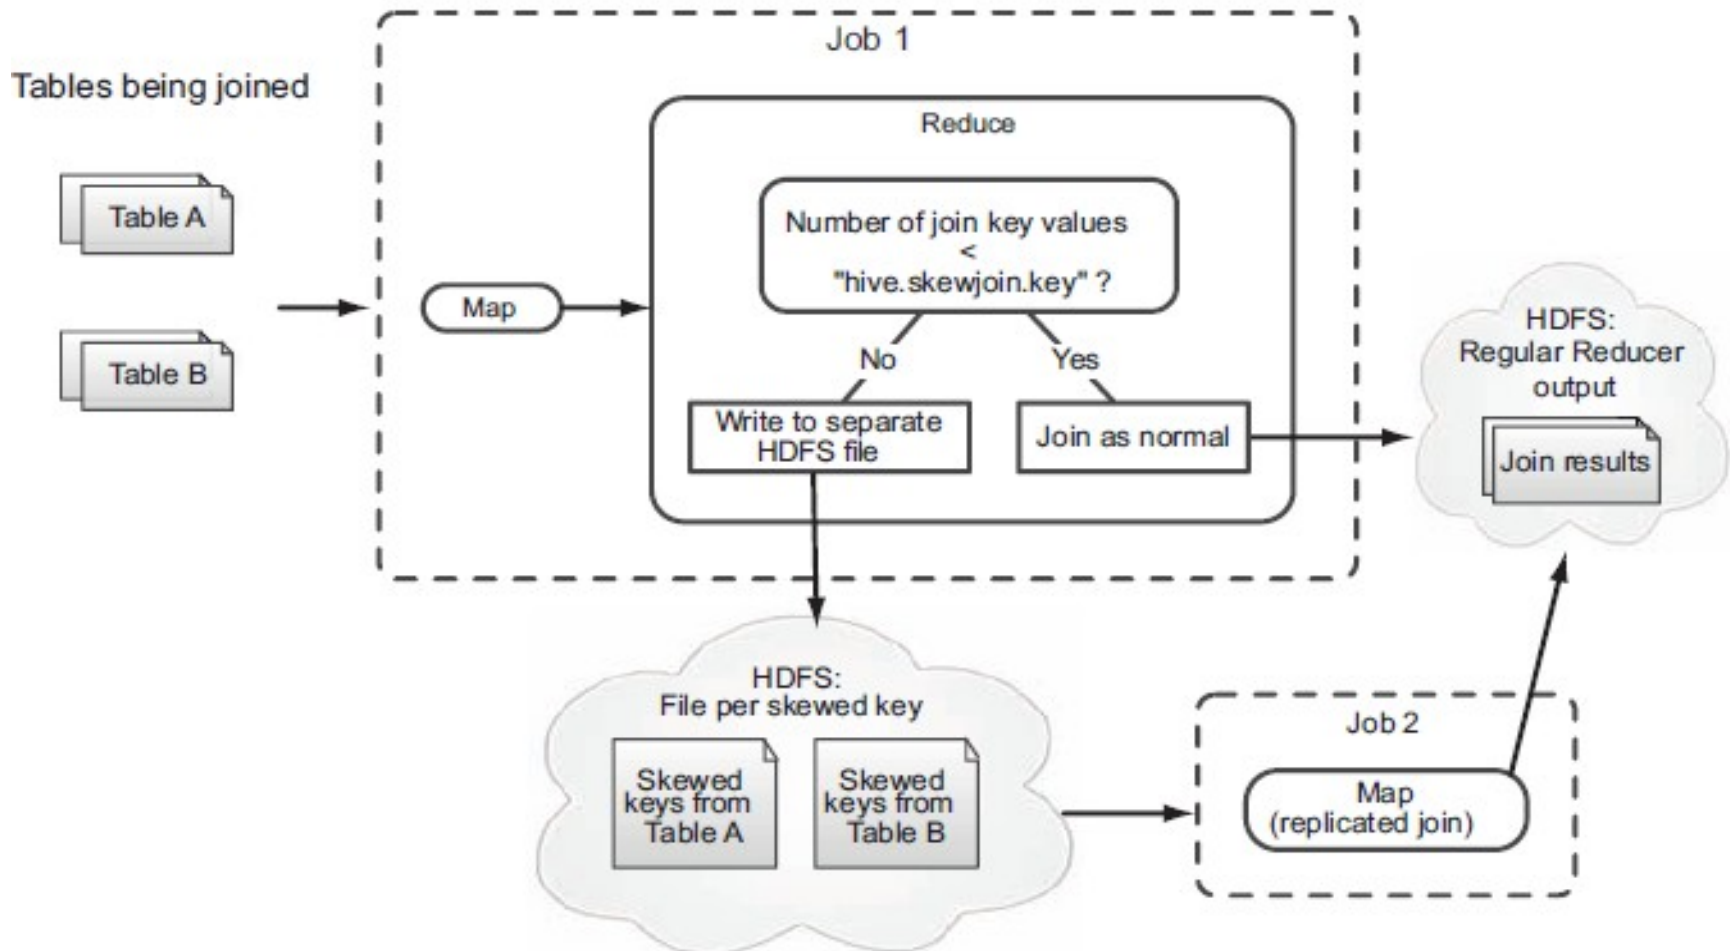
```
hive> SELECT  viewed_movies.movie  
FROM viewed_movies  
LEFT SEMI JOIN movie_categories  
ON viewed_movies.movie = movie_categories.title;
```

```
127Hours  
BlackSwan  
TheFighter
```

In a semi-join you can't have the result contain any fields from the right-hand side of the join, which in your case is the `movie_categories` table.



# Hive:data analytics:example-1:skew-joins



# Hive:data

## analytics:example-1:skew-joins

```
hive> SET hive.optimize.skewjoin = true;
```

← Tell Hive to optimize joins where it sees skewed data.

```
hive> SET hive.skewjoin.key = 100000;
```

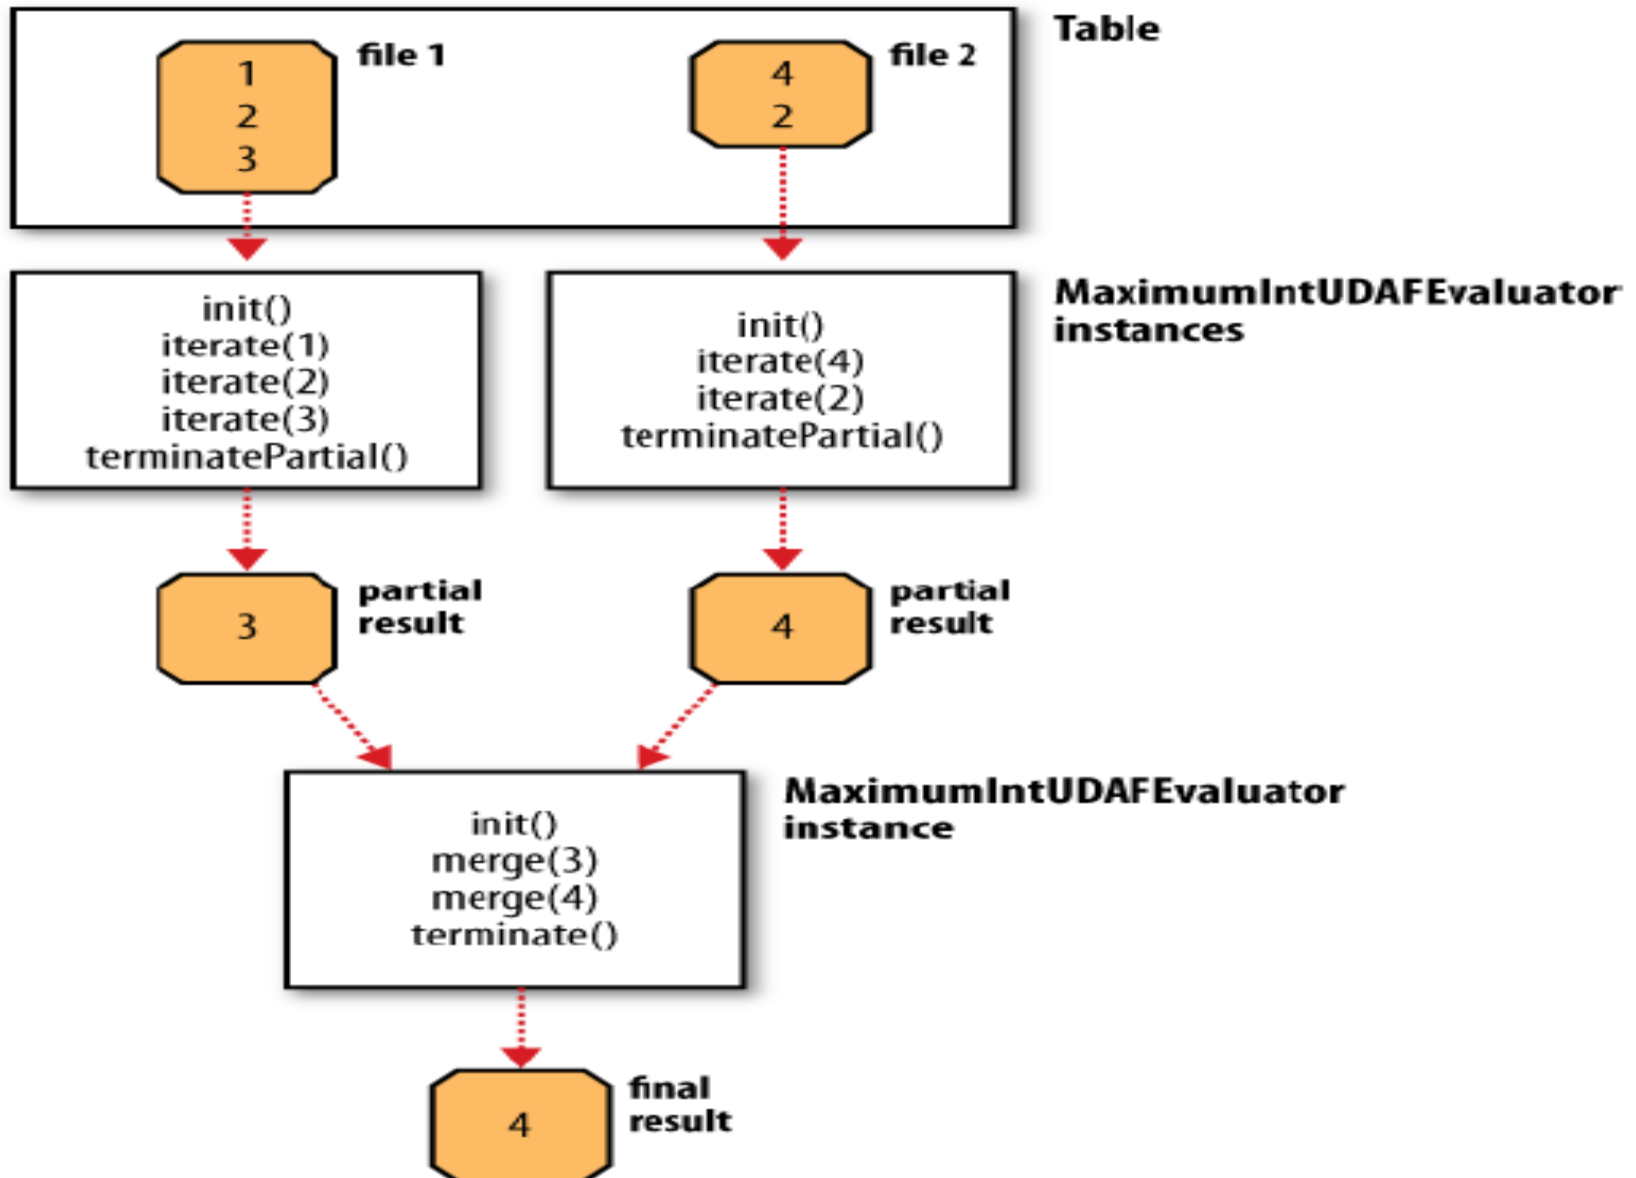
← Sets the threshold beyond which a key is considered to be skewed.



# Hive:UDAFs

```
public class Maximum extends UDAF {  
  
    public static class MaximumIntUDAFEvaluator implements UDAFEvaluator {  
  
        private IntWritable result;  
  
        public void init() {  
            result = null;  
        }  
        public boolean iterate(IntWritable value) {  
            if (value == null) {  
                return true;  
            }  
            if (result == null) {  
                result = new IntWritable(value.get());  
            } else {  
                result.set(Math.max(result.get(), value.get()));  
            }  
            return true;  
        }  
  
        public IntWritable terminatePartial() {  
            return result;  
        }  
  
        public boolean merge(IntWritable other) {  
            return iterate(other);  
        }  
  
        public IntWritable terminate() {  
            return result;  
        }  
    }  
}
```

# Hive:UDAFs



# Hive:OrderBy vs SortBy

```
hive> SELECT movie_categories.category, count(1) AS cnt
FROM viewed_movies
JOIN movie_categories ON
    (viewed_movies.movie = movie_categories.title)
WHERE viewed_movies.country = "RU"
GROUP BY movie_categories.category
ORDER BY cnt DESC;
```

Search for the top categories in Russia.

Your table is partitioned by country,  
which speeds up your query because

only files contained in these partitions will be loaded.

Order the results by the number of movies in each category,  
so the most popular category is the first result.

# Hive:OrderBy vs SortBy:Explain

ABSTRACT SYNTAX TREE:

```
(TOK_QUERY (TOK_FROM (TOK_JOIN (TOK_TABREF  
(TOK_TABNAME viewed_movies)) (TOK_TABREF  
(TOK_TABNAME movie_categories)) (= (.  
...
```

STAGE DEPENDENCIES:

Stage-1 is a root stage

Stage-2 depends on stages: Stage-1

Stage-3 depends on stages: Stage-2

Stage-0 is a root stage

## Stage: Stage-1

### Map Reduce

Alias -> Map Operator Tree:

movie\_categories

TableScan

alias: movie\_categories

This tells you that the movie\_categories table is one of the input tables for the job.

Reduce Output Operator

key expressions:

expr: title

type: string

This tells you that the map is emitting an output key/value tuple where the key is the movie title.

sort order: +

Map-reduce partition columns:

expr: title

type: string

The title, which is the output key, is being used for partitioning.

tag: 1

value expressions:

expr: category

type: string

The output value is the movie category.

viewed\_movies

This section shows details for the viewed\_movies table. The output key/value tuples are the movie title and country, respectively.

TableScan

alias: viewed\_movies

Reduce Output Operator

key expressions:

expr: movie

type: string

sort order: +

Map-reduce partition columns:

expr: movie

type: string

tag: 0

value expressions:

expr: country

type: string

Reduce Operator Tree: The first operation on the reduce side is a join, and by looking down a couple of lines you can see that it's an inner join.

Join Operator

condition map:

Inner Join 0 to 1

condition expressions:

0 {VALUE.\_col3}

1 {VALUE.\_col1}

handleSkewJoin: false

outputColumnNames: \_col3, \_col7

Filter Operator

predicate:

expr: (\_col3 = 'RU')

type: boolean

Select Operator

expressions:

expr: \_col7

type: string

outputColumnNames: \_col7

Group By Operator

aggregations:

expr: count()

bucketGroup: false

keys:

expr: \_col7

type: string

mode: hash

outputColumnNames: \_col0, \_col1

File Output Operator

compressed: true

GlobalTableId: 0

table:

input format: org.apache.hadoop.mapred.

SequenceFileInputFormat

output format: org.apache.hadoop.hive ql.io.

HiveSequenceFileOutputFormat

This operator is filtering results based on your criteria that only results from Russia be included.

The select operator does nothing other than pass through data to the next operator.

This group is a localized group within all the rows for a given join key. You're grouping on the movie category, so the output of this group is the category name and a count of the number of rows with the same category.

The last step indicates that the output is being written to file and the output format used to perform the write.